


1st: Method: GLM

In cleaned data, there are 16 numerical predictors and 6 categorical predictors.

1. Use correlation plot of numerical predictors to seek interactions with $\text{abs}(\text{correlation}) > 0.5$
2. Build GLM based only on numerical predictors and drop insignificant terms
3. Combine numerical terms and their interaction with each categorical predictor in different GLM, e.g. numerical terms+numerical terms:some categorical, and use drop-in deviance test to determine whether a categorical predictor is useful
4. Relevel categorical predictor and combine all interactions into the full model
5. Forward selection of 20 terms based on AIC starting from intercept-only model.
6. Determine the classifier to obtain categorical fitted value.



❖ Prediction accuracy = 0.804

```
Call:
glm(formula = status ~ int_rate + dti + tot_cur_bal + emp_length +
    total_acc + annual_inc + delinq_2yrs + open_acc + percent_bc_gt_75 +
    term + inq_last_6mths + verification_status + pct_tl_nvr_dlq +
    installment + loan_amnt:annual_inc + tot_cur_bal:annual_inc +
    loan_amnt:installment, family = binomial(), data = data[,
    c(-10)])
```

Coefficients:


	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	3.636e+00	1.422e-01	25.571	< 2e-16	***
int_rate	-9.620e+00	2.598e-01	-37.035	< 2e-16	***
dti	-2.676e-02	1.295e-03	-20.669	< 2e-16	***
tot_cur_bal	1.431e-06	1.119e-07	12.784	< 2e-16	***
emp_length1 year	9.417e-02	4.769e-02	1.975	0.048291	*
emp_length10+ years	2.298e-01	3.573e-02	6.432	1.26e-10	***
emp_length2 years	1.319e-01	4.377e-02	3.012	0.002592	**
emp_length3 years	1.354e-01	4.534e-02	2.986	0.002827	**
emp_length4 years	1.493e-01	4.980e-02	2.997	0.002723	**
emp_length5 years	7.409e-02	4.909e-02	1.509	0.131251	
emp_length6 years	1.114e-01	5.056e-02	2.204	0.027544	*
emp_length7 years	1.535e-01	4.985e-02	3.079	0.002078	**
emp_length8 years	1.077e-01	5.102e-02	2.111	0.034758	*
emp_length9 years	9.613e-02	5.451e-02	1.764	0.077791	.
emp_lengthn/a	-2.269e-01	5.014e-02	-4.525	6.05e-06	***
total_acc	1.273e-02	1.117e-03	11.401	< 2e-16	***
annual_inc	6.977e-06	6.665e-07	10.468	< 2e-16	***
delinq_2yrs	-9.688e-02	1.098e-02	-8.822	< 2e-16	***
open_acc	-1.907e-02	2.501e-03	-7.624	2.47e-14	***
percent_bc_gt_75	-2.571e-03	2.833e-04	-9.078	< 2e-16	***
term 60 months	-2.262e-01	2.415e-02	-9.364	< 2e-16	***
inq_last_6mths	-5.049e-02	8.465e-03	-5.965	2.44e-09	***
verification_statusVerified	-6.117e-02	2.325e-02	-2.631	0.008506	**
pct_tl_nvr_dlq	-4.830e-03	1.313e-03	-3.680	0.000233	***
installment	-7.746e-04	1.408e-04	-5.501	3.77e-08	***
annual_inc:loan_amnt	-1.690e-10	2.939e-11	-5.753	8.79e-09	***
tot_cur_bal:annual_inc	-2.496e-12	6.093e-13	-4.096	4.20e-05	***
installment:loan_amnt	1.916e-08	4.508e-09	4.251	2.13e-05	***

Some interpretation of coefficients:

Annual income has positive coefficient, meaning higher income indicates larger capacity to repay;


Lower interest rate indicates larger capacity to repay;

Employment length:




2nd Method: RandomForest

- ❖ 5 fold Cross-Validations for feature selection.




- ❖ Check importance plot and remove non-significant predictors.
- ❖ Rebuild model with 12 predictors

int_rate
annual_inc
tot_cur_bal
avg_cur_bal
installment
loan_amnt
open_acc
dti
total_acc
percent_bc_gt_75
mort_acc
home_ownership
pct_tl_nvr_dlq
num_accts_ever_120_pd
num_bc_sats
verification_status
emp_length
delinq_2yrs
inq_last_6mths




output3

dti
emp_length
int_rate
installment
annual_inc
avg_cur_bal
tot_cur_bal
loan_amnt
total_acc
open_acc
percent_bc_gt_75
num_bc_sats
pct_tl_nvr_dlq
inq_last_6mths
verification_status
num_accts_ever_120_pd
delinq_2yrs
home_ownership
mort_acc



❖ Tune RandomForest for the optimal mtry parameter



Optimal Parameter:

- ❖ $ntree = 500$, number of trees to grow
- ❖ $mtry = 3$, number of variables randomly sampled as candidates at each split.
- ❖ Prediction accuracy = 0.804

3rd Method: SVM with RBF kernel

- ❖ We normalized the dataset (training and testing) for SVM. SVMs assume that the data it works with is in a standard range (0 to 1 roughly). So the normalization of feature vectors prior to feeding them to the SVM is very important. We want to make sure that for each dimension, the values are scaled to lie roughly within this range.
- ❖ We selected optimal parameters (the cost and gamma) by 10-folds CV.
 - Parameter tuning of 'svm':
 - sampling method: 10-fold cross validation
 - best parameters:
cost gamma
0.1 0.1
 - best performance: 0.1945
- ❖ Prediction accuracy = 0.8055

- ❖ Train the model with optimal parameters(cost = 0.1, gamma = 0.1) with 10-folds CV.

```
Call:
svm.default(x = newtrain, y = train_label,
            kernel = "radial", gamma = 0.1, cost = 0.1,
            cross = 10)

Parameters:
  SVM-Type:  C-classification
 SVM-Kernel: radial
      cost:  0.1
      gamma: 0.1

Number of Support Vectors: 32418
( 12921 19497 )

Number of Classes: 2

Levels:
0 1



10-fold cross-validation on training data:

Total Accuracy: 80.20983
Single Accuracies:
79.9663 80.30326 80.10415 80.73212 79.62935 79.69061 80.96186 79.95099 80.31858 80.
44111
```

4th Method: AdaBoosting and Bagging

- ❖ adaBoosting with total 50 iterations for which the number of trees to use. Additionally, a bootstrap sample of the training set is drawn using the weights for each observation on that iteration.
- ❖ adaBoosting Prediction accuracy = 0.797
- ❖ Bagging Prediction accuracy = 0.804

Variables relative importance



5th Method: Extreme Gradient Boosting

Optimal Parameters for Tree/Linear Booster:

- ❖ eta = 0.001 (control the learning rate), which used to prevent overfitting by making the boosting process more conservative. (for tree only)
- ❖ max.depth = 4, maximum depth of a tree (for tree only)
- ❖ subsample = 0.8, randomly collected 80% of the data instances to grow trees and this will prevent overfitting.

Tree Booster converged at 7th iteration

```
[101] train-error:0.190062 test-error:0.186250
[102] train-error:0.190062 test-error:0.186250
[103] train-error:0.190062 test-error:0.186250
[104] train-error:0.190062 test-error:0.186250
[105] train-error:0.190062 test-error:0.186250
[106] train-error:0.190062 test-error:0.186250
Stopping. Best iteration: 7
```

Linear Booster converged at 2th iteration

```
[95] train-error:0.185937 test-error:0.188500
[96] train-error:0.186125 test-error:0.188000
[97] train-error:0.185875 test-error:0.188500
[98] train-error:0.186000 test-error:0.188250
[99] train-error:0.185937 test-error:0.188250
[100] train-error:0.185875 test-error:0.188500
[101] train-error:0.185875 test-error:0.188250
Stopping. Best iteration: 2
```

Final Result

	AccuracyRate	Method
1	0.80250	RandomForest
2	0.80550	SVM_RBF
3	0.79700	adaBoosting
4	0.80400	Bagging
5	0.81175	xgBoosting_tree
6	0.81375	xgBoosting_linear

